

# Reinforcement Learning

Studiengang: Data Science and Business Analytics


Vortragender: Sebastian Eresheim

Ort, Datum: St. Pölten, 05.12.2021

# Organizational

Lecture is based on:

- *Reinforcement Learning: An Introduction*  
Richard Sutton, Andrew Barto  
[www.incompleteideas.net/book/the-book.html](http://www.incompleteideas.net/book/the-book.html)



## Machine Learning

Supervised Learning

Unsupervised Learning

Reinforcement Learning

Reinforcement Learning is motivated by learning:

- like a human
- via interaction (Trial and Error)
- „learning what to do!“ not „how to do it“.


# Overview

In reinforcement learning an agent learns how to map actions to observations in order to maximise a numerical reward, which it receives from an environment.


# Overview

In reinforcement learning an **agent** learns how to map actions to observations in order to maximise a numerical reward, which it receives from an **environment**.

# Preliminaries

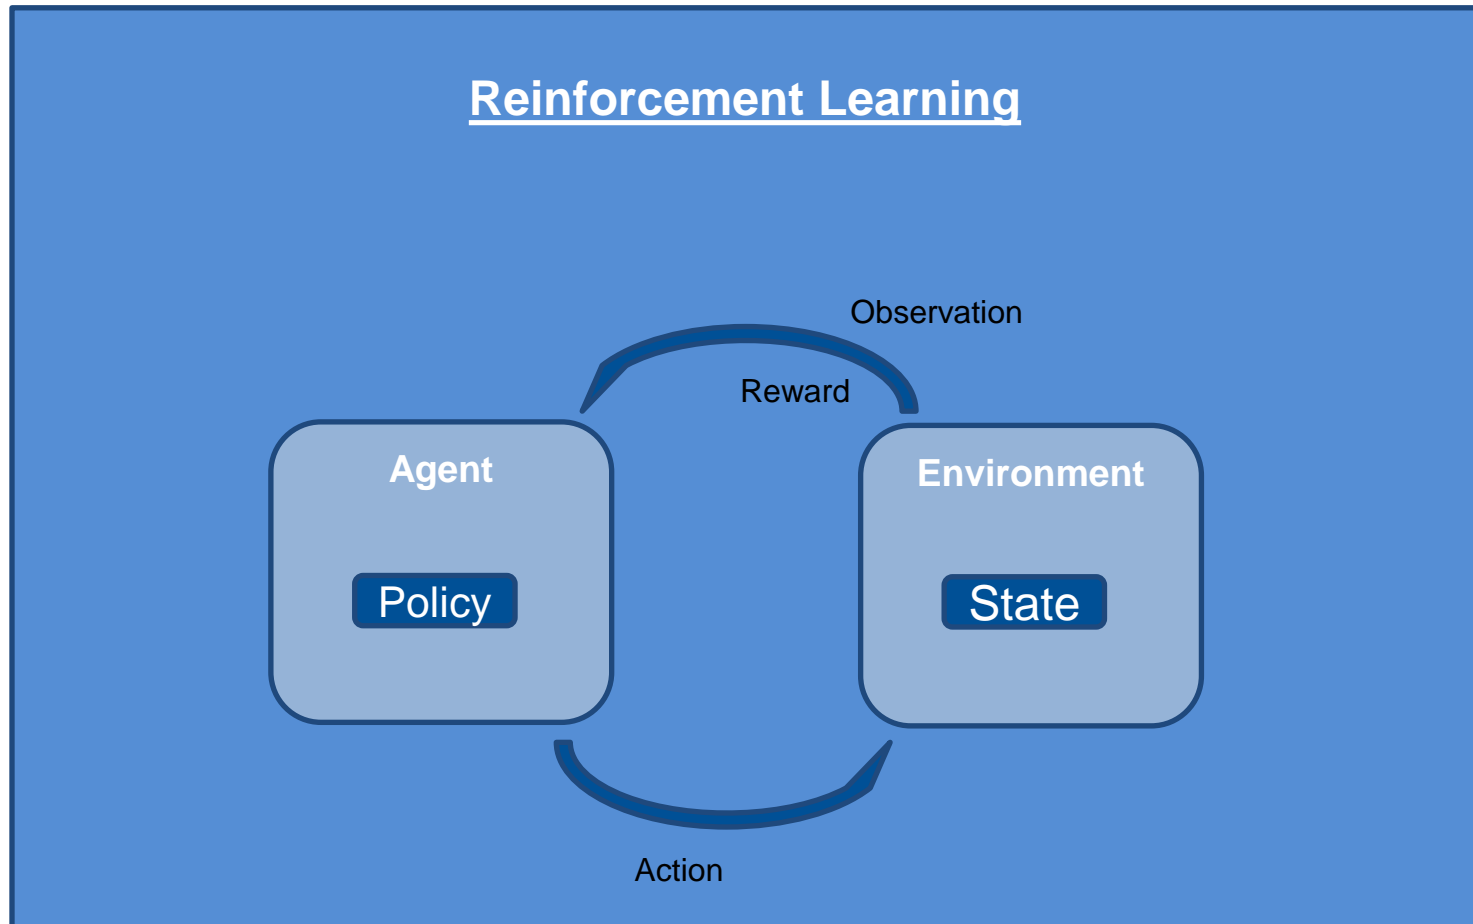


# Preliminaries

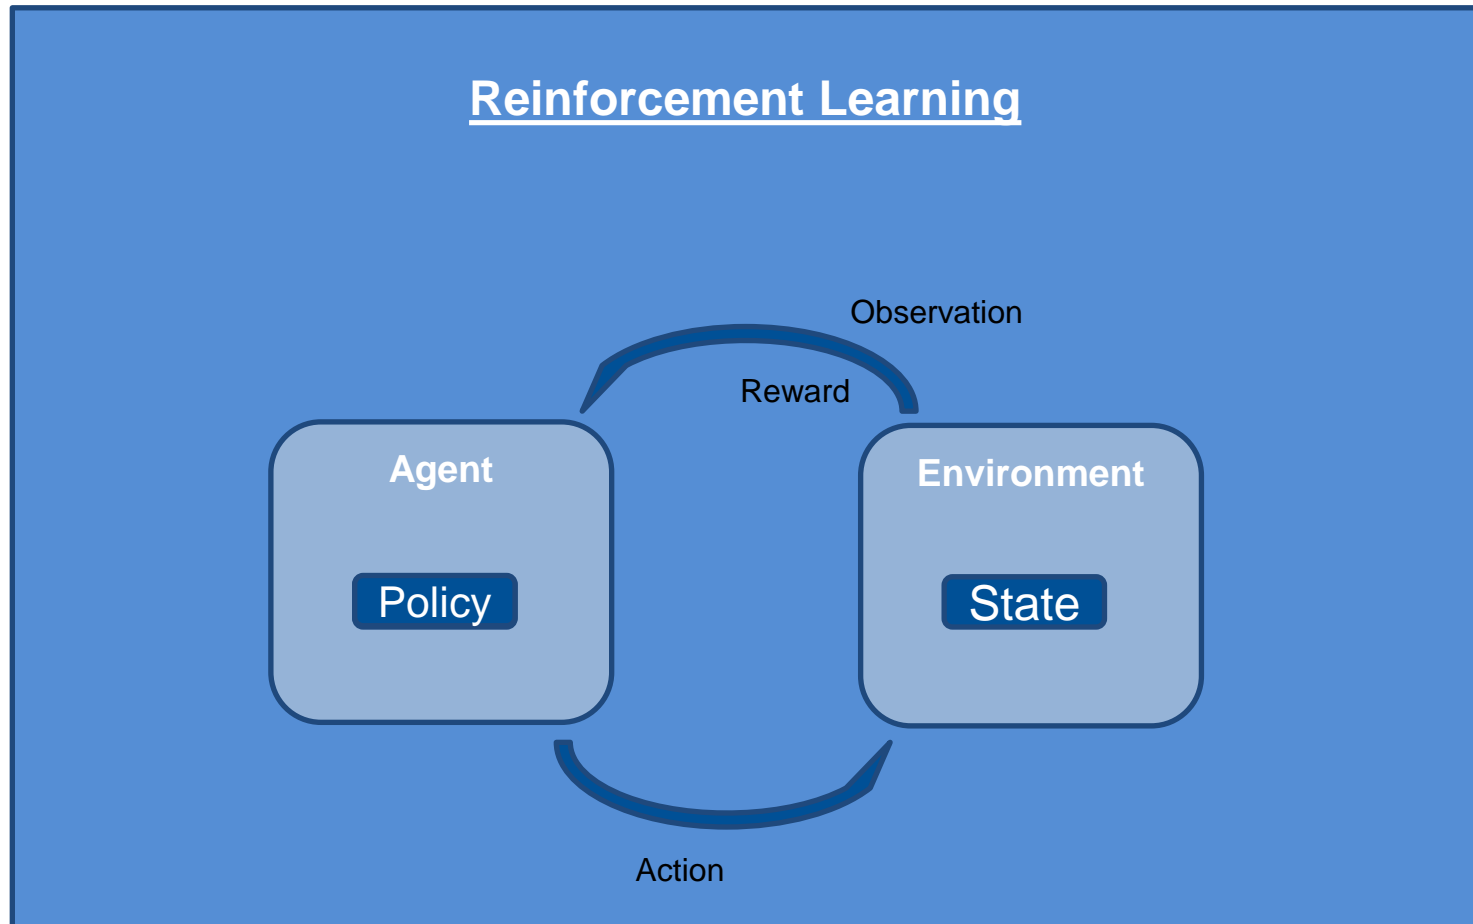




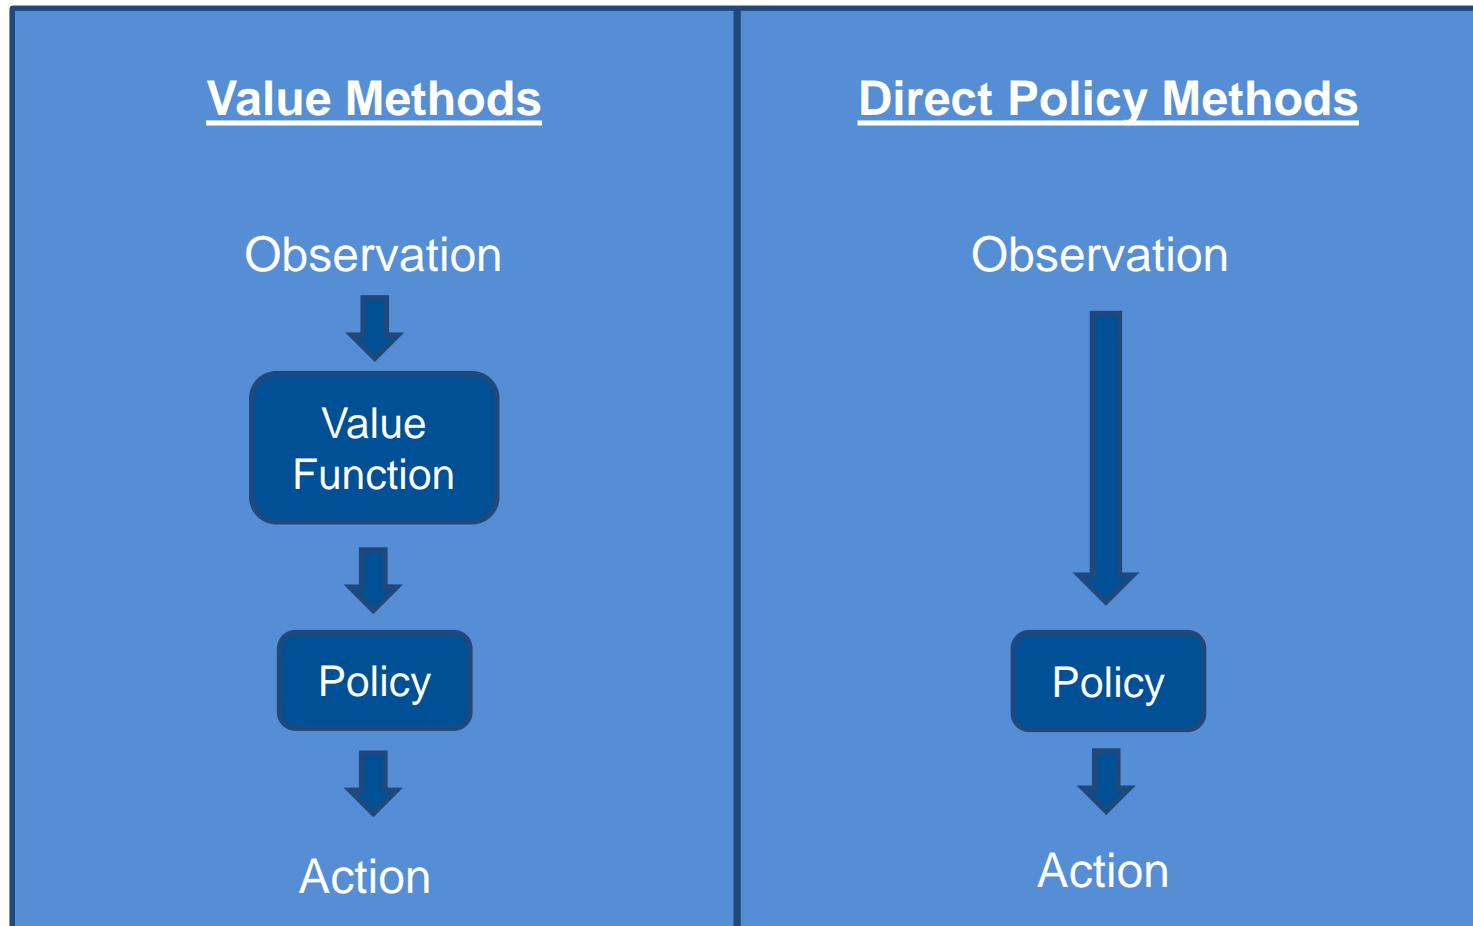
# Preliminaries



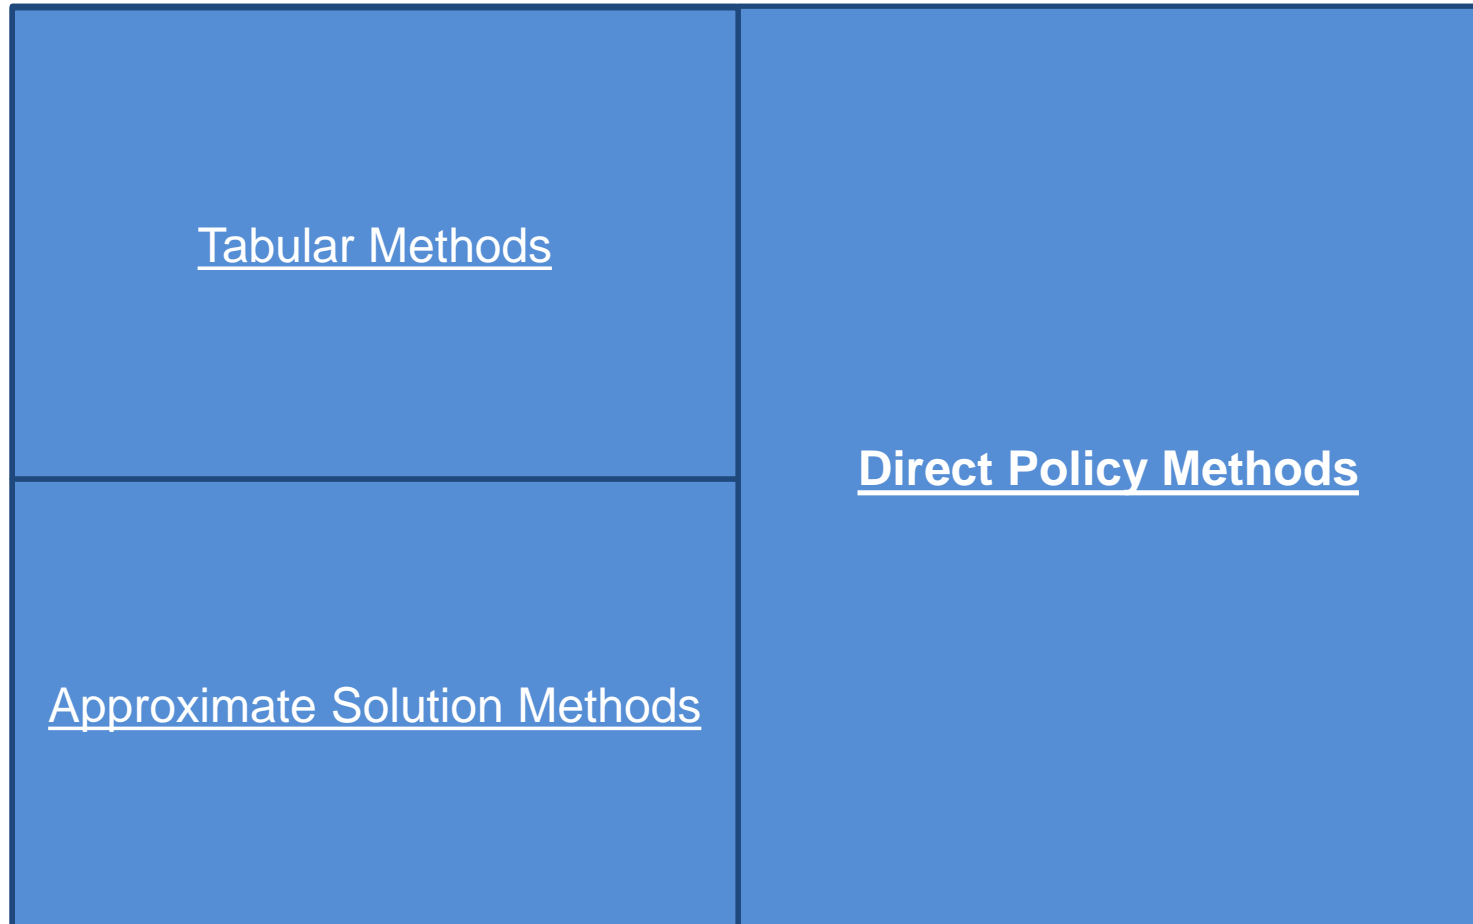
# Preliminaries



# Preliminaries



# Preliminaries



# Overview - Problems


Exploitation

vs

Exploration


# Overview - Problems

## Delayed Reward




# Overview - Problems

## Credit Assignment Problem



# Overview - Problems


## Non-Stationarity





# Overview - Problems

## Partial Observability




# Overview - Problems

## Multi-Agent



Image: <https://deepmind.com/blog/article/capture-the-flag-science>

# Milestones of RL

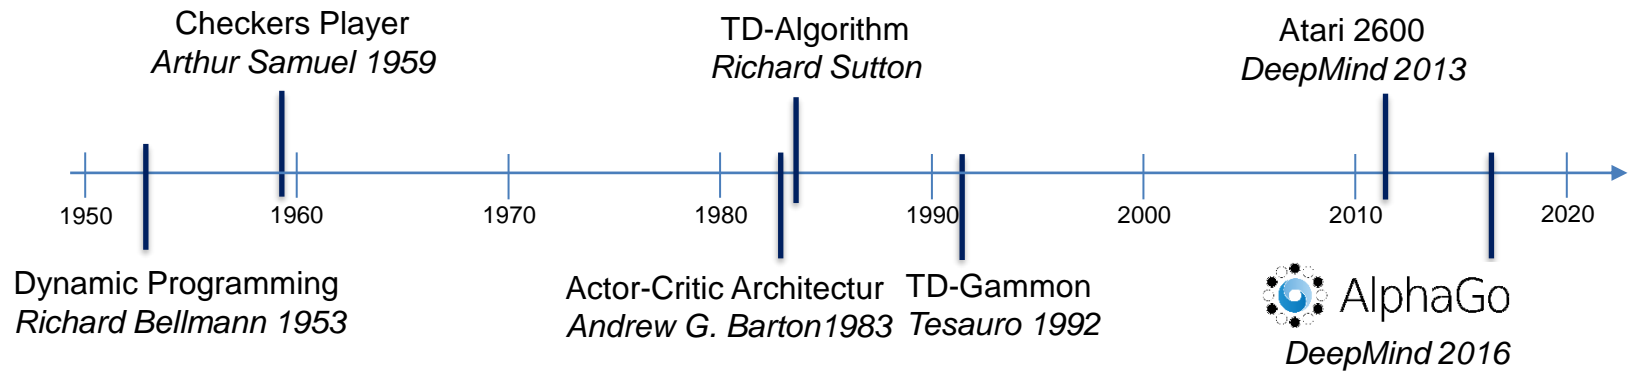


# Atari Games

**Starting out - 10 minutes of training**

**The algorithm tries to hit the ball back, but  
it is yet too clumsy to manage.**

# Milestones of RL




# AlphaGo – AlphaZero

/informatik & security



Image: [https://cdn-images-1.medium.com/max/1600/0\\*J2HH4TDqGCuKbufS.png](https://cdn-images-1.medium.com/max/1600/0*J2HH4TDqGCuKbufS.png)

# Milestones of RL





# Open AI 5




Source: Youtube [https://www.youtube.com/watch?v=yBEidvm\\_tZQ](https://www.youtube.com/watch?v=yBEidvm_tZQ)

Reinforcement Learning| Wintersemester 2019 | Sebastian Eresheim



# Milestones of RL



# AlphaStar

/informatik & security


**fh** st. pölten



Source: Youtube <https://www.youtube.com/watch?v=cUTMhmVh1qs>

Reinforcement Learning| Wintersemester 2019 | Sebastian Eresheim

# Milestones of RL



# Multi-armed Bandits

Exploitation

vs

Exploration



# INTERACTIVE PART I

# Multi-armed Bandits

**n** choices of **k** different actions

After each choice: numerical reward from stationary distribution

Objective: maximize total reward over  $n$  choices

# Multi-armed Bandits

$A_t$ ... Action at timestep  $t$

$R_t$ ... Reward at timestep  $t$

optimal action-value function:

(for the multi-armed bandit problem)

$$q_*(a) := \mathbb{E}[R_t | A_t = a]$$

We usually don't know that value!

$q_t(a)$ ... Estimate of  $q_*(a)$  at timestep  $t$



# INTERACTIVE PART II



# Multi-armed Bandits

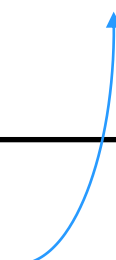
Estimator of action-value function:

$$q_t(a) := \frac{\text{sum of rewards when } a \text{ taken before } t}{\text{number of times } a \text{ taken before } t} = \frac{\sum_{i=1}^{t-1} R_i \cdot \mathbb{1}_{[A_i=a]}}{\sum_{i=1}^{t-1} \mathbb{1}_{[A_i=a]}}$$

Indicator function:

$$\mathbb{1}_{[A_i=a]}(a) := \begin{cases} 1 & \text{if } A_i = a \\ 0 & \text{if } A_i \neq a \end{cases}$$

complicated way of  
saying:  
average of rewards



# Action Selection Methods

Random action selection method:

Randomly select one of the possible actions.

# Action Selection Methods

Graph of a probability mass function (pmf) of the distribution:



No exploitation - Only exploration

## Greedy action selection method:

Always select the action with the highest action-value according to the action-value function.

$$A_t = \operatorname{argmax}_a q_t(a)$$

# Action Selection Methods

Graph of a probability mass function (pmf) of the distribution:



Only exploitation - no exploration

$\varepsilon$ -Greedy action selection method:

In most  $(1-\varepsilon)$  cases select the best action and in a small amount of cases select a random action.

$$A_t = \begin{cases} \operatorname{argmax}_a q_t(a) & \text{in } (100-\varepsilon) \% \text{ cases} \\ \text{random action} & \text{in } \varepsilon \% \text{ cases} \end{cases}$$

Mostly exploitation - small exploration



# Action Selection Methods

Graph of a probability mass function (pmf) of the distribution:



Mostly exploitation - small exploration

# Multi-armed Bandits





# Incremental Implementation

Simplified: only one action

$$q_n = \frac{R_1 + R_2 + \dots + R_n}{n}$$

$$\begin{aligned} q_n &= \frac{1}{n} \sum_{i=1}^n R_i \\ &= \frac{1}{n} (R_n + \sum_{i=1}^{n-1} R_i) \\ &= \frac{1}{n} (R_n + (n-1) \frac{1}{n-1} \sum_{i=1}^{n-1} R_i) \\ &= \frac{1}{n} (R_n + (n-1) Q_{n-1}) \\ &= \frac{1}{n} (R_n + n Q_{n-1} - Q_{n-1}) \\ &= Q_{n-1} + \frac{1}{n} (R_n - Q_{n-1}) \end{aligned}$$

# Incremental Implementation

/informatik & security



$$q_{n+1} = q_n + \frac{1}{n+1} (R_n - q_n)$$

$$NewEstimate \leftarrow OldEstimate + StepSize \underbrace{[Target - OldEstimate]}_{Error}$$

# Incremental Implementation


/informatik & security

$$NewEstimate \leftarrow OldEstimate + StepSize \underbrace{[Target - OldEstimate]}_{\text{Error}}$$

$$R_1 = 1,93 \quad Q_1 = 0 + \frac{1}{1} [1,93 - 0] = 1,93$$

$$R_2 = 2,35 \quad Q_2 = 1,93 + \frac{1}{2} [2,35 - 1,93] = 2,14$$

$$R_3 = 3,32 \quad Q_3 = 2,14 + \frac{1}{3} [3,32 - 2,14] = 2,53$$




# Incremental Implementation

$$NewEstimate \leftarrow OldEstimate + StepSize \underbrace{[Target - OldEstimate]}_{\text{Error}}$$

$$R_4 = 2,43 \quad Q_4 = 2,53 + \frac{1}{4}[2,43 - 2,53] = 2,505$$

$$R_5 = 3,28 \quad Q_5 = 2,505 + \frac{1}{5}[3,28 - 2,505] = 2,66$$

$$R_6 = 3,91 \quad Q_6 = 2,66 + \frac{1}{6}[3,91 - 2,66] = 2,868$$




# Incremental Implementation

$$NewEstimate \leftarrow OldEstimate + StepSize \underbrace{[Target - OldEstimate]}_{\text{Error}}$$

$$R_7 = 2,34 \quad Q_7 = 2,868 + \frac{1}{7}[2,34 - 2,868] = 2,793$$

$$R_8 = 3,42 \quad Q_8 = 2,793 + \frac{1}{8}[3,42 - 2,793] = 2,871$$


$$R_9 = 2,64 \quad Q_9 = 2,871 + \frac{1}{9}[2,64 - 2,871] = 2,845$$




# Incremental Implementation

/informatik & security

$$NewEstimate \leftarrow OldEstimate + StepSize \underbrace{[Target - OldEstimate]}_{Error}$$



# Optimistic initial values



# Nonstationary Problem


Constant step-size Parameter  $\alpha \in (0,1]$        $Q_{n+1} := Q_n + \alpha(R_n - Q_n)$

$$\begin{aligned} Q_{n+1} &= Q_n + \alpha(R_n - Q_n) \\ &= \alpha R_n + (1 - \alpha)Q_n \\ &= \alpha R_n + (1 - \alpha)[\alpha R_{n-1} + (1 - \alpha)Q_{n-1}] \\ &= \alpha R_n + (1 - \alpha) \alpha R_{n-1} + (1 - \alpha)^2 Q_{n-1} \\ &= \alpha R_n + (1 - \alpha) \alpha R_{n-1} + (1 - \alpha)^2 \alpha R_{n-2} + \dots + (1 - \alpha)^{n-1} \alpha R_1 + (1 - \alpha)^n Q_1 \\ &= (1 - \alpha)^n Q_1 + \sum_{i=1}^n \alpha (1 - \alpha)^{n-i} R_i \end{aligned}$$



# Contextual Bandits

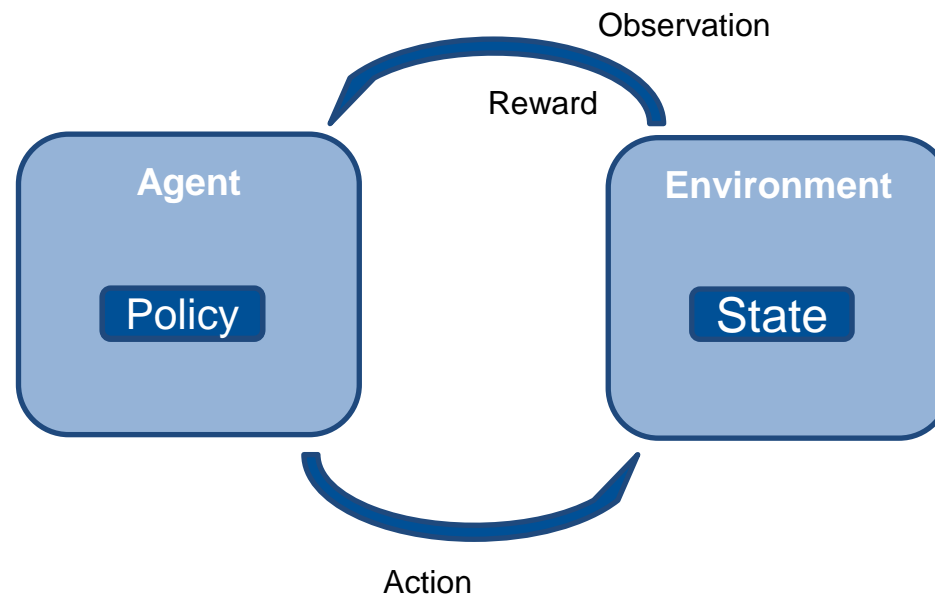
- Synonym for the full reinforcement learning problem
- Multiple situations in contrast to only one
- Action also affects the next situation not only the reward (Return)




# Multi-armed Bandits

- In which subfields can RL be divided into? What are the major differences?
- What is the multi-armed bandit problem? What is the difference to the full RL problem?
- What is an action-value function? How is its estimator defined?
- What is the incremental notation of an action-value function estimator?
- Which action selection methods exist?

# MDP




# MDP



# MDP – Example

Tic-Tac-Toe



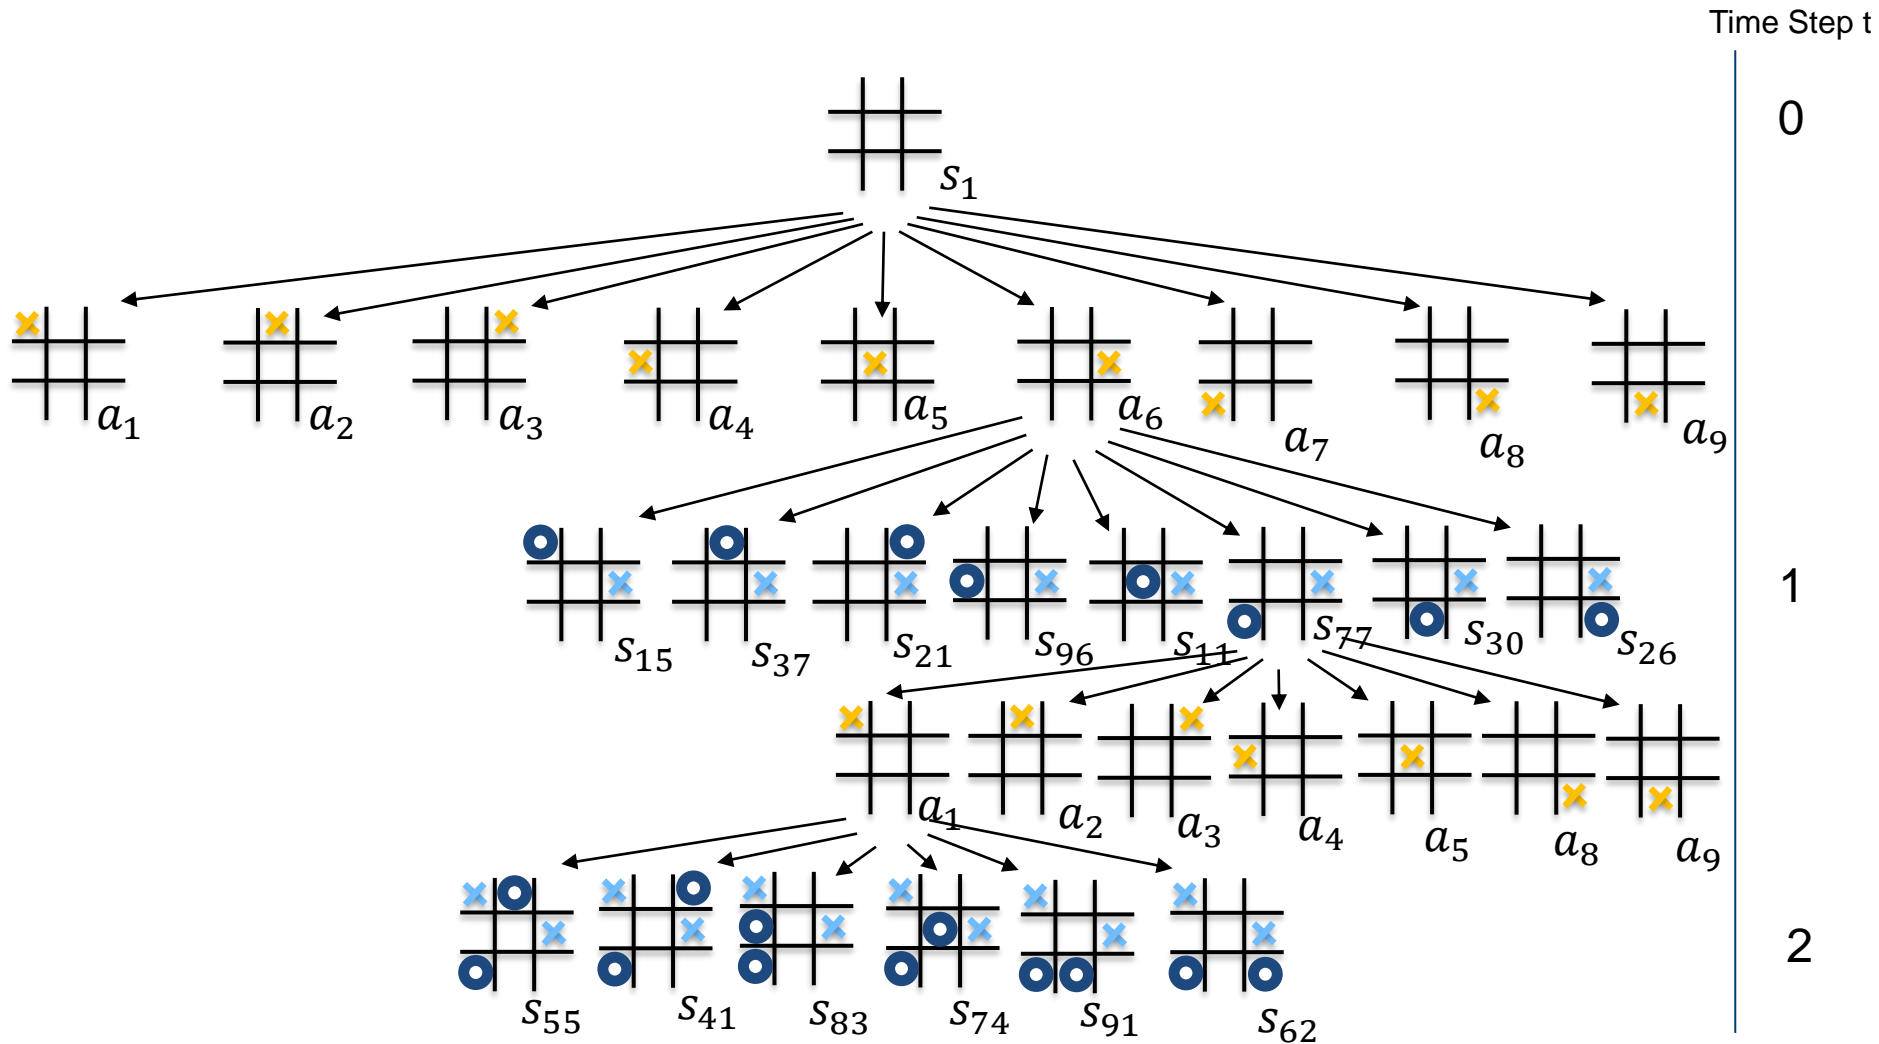
# MDP – Example

State Space:  $\mathcal{S} = \{ s_1, s_2, s_3, s_4, s_5, \dots \}$

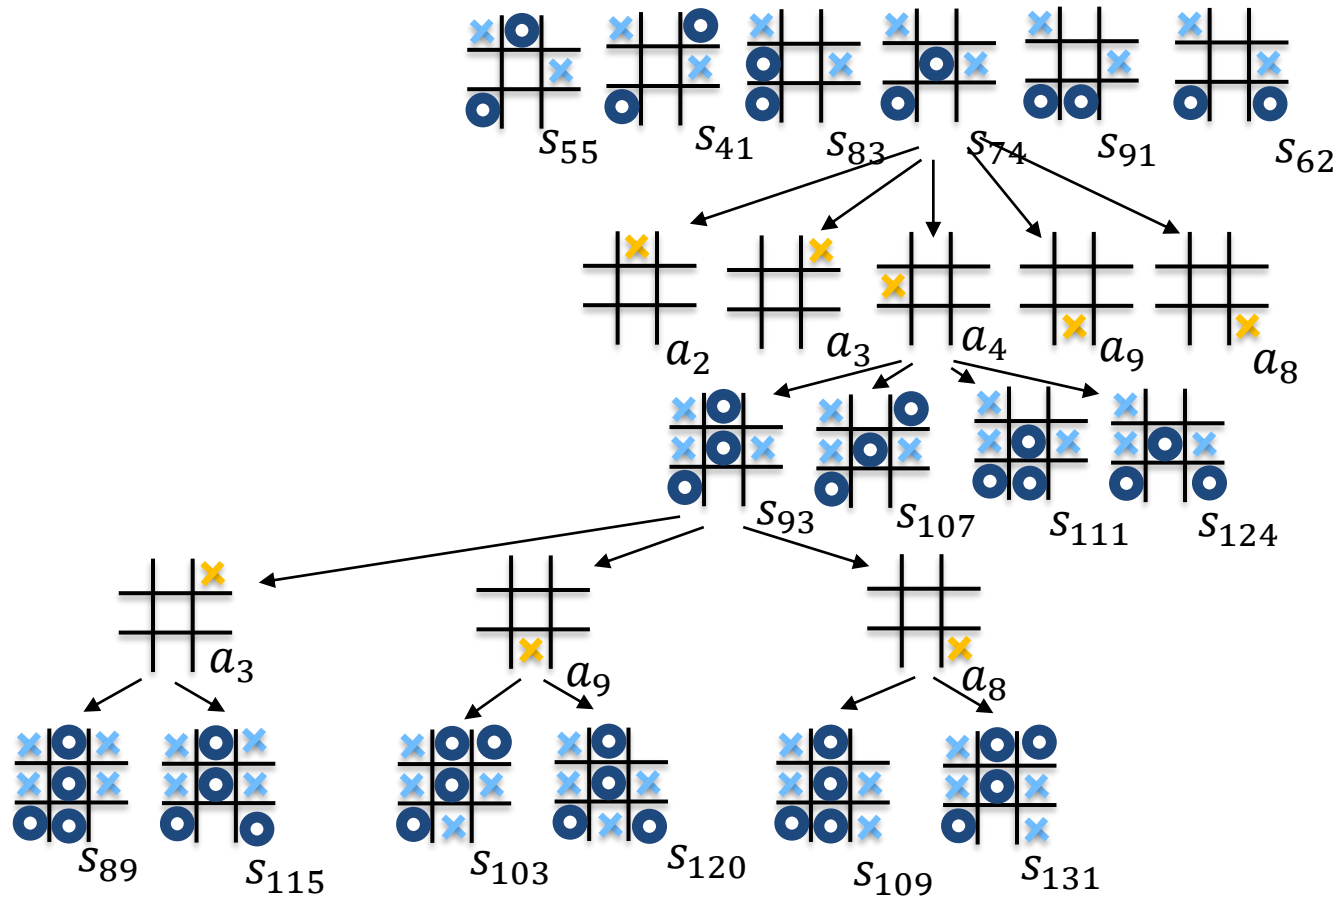
Action Space:  $\mathcal{A} = \{ a_1, a_2, a_3, a_4, a_5, a_6, a_7, a_8, a_9 \}$

Reward Function: 
$$\begin{cases} +1 & \text{if } \begin{matrix} \text{blue 'x's at (1,1), (1,2), (1,3)} \\ \text{blue 'x's at (1,1), (2,1), (2,2), (2,3)} \\ \text{blue 'x's at (1,1), (2,2), (3,3)} \end{matrix} \text{, or} \\ -1 & \text{if } \begin{matrix} \text{blue circles at (1,1), (1,2), (1,3)} \\ \text{blue circles at (1,1), (2,1), (2,2), (2,3)} \\ \text{blue circles at (1,1), (2,2), (3,3)} \end{matrix} \text{, or} \\ 0 & \text{else} \end{cases}$$

# MDP – Example



# MDP – Example



Time Step t

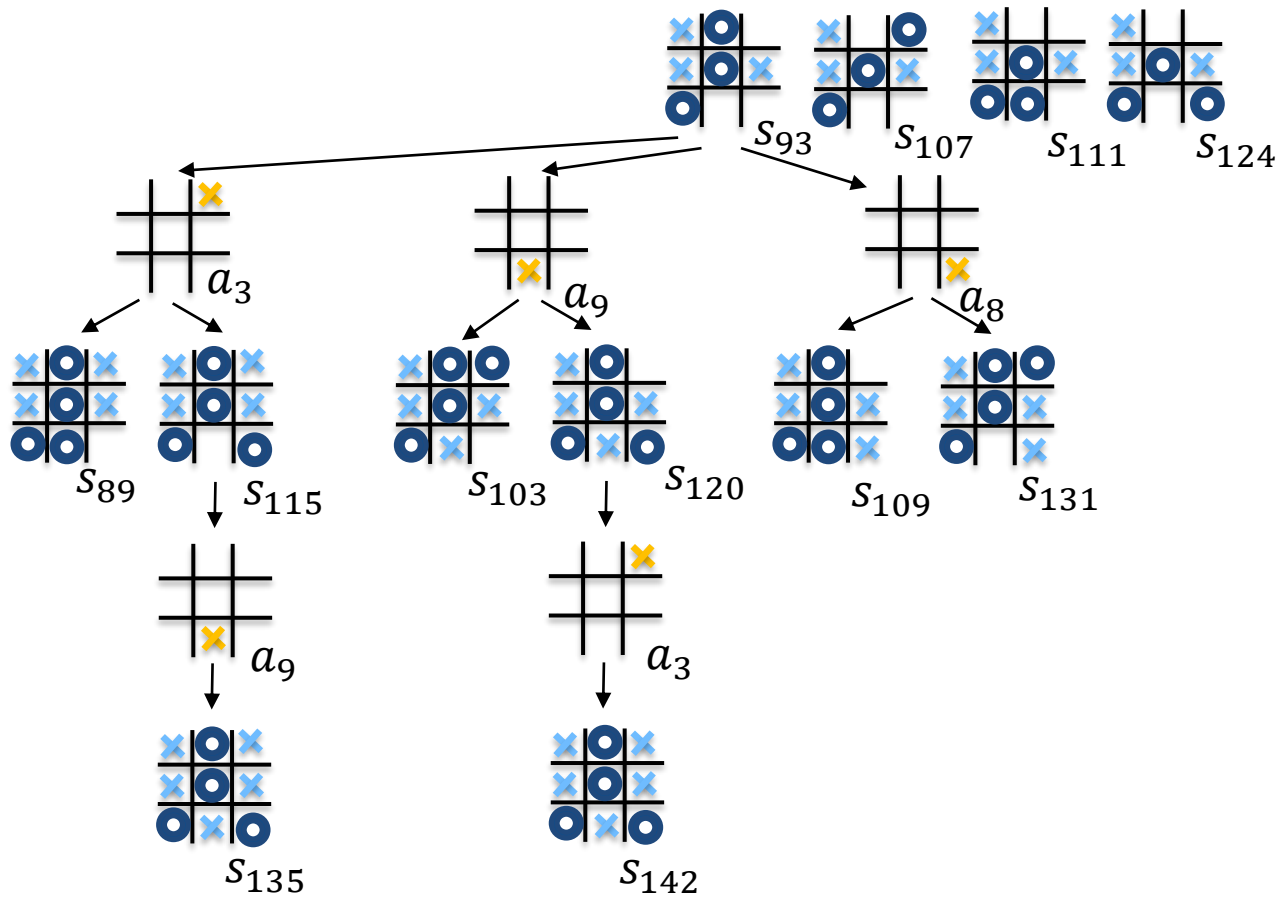
2

3

4



# MDP – Example




Time Step  $t$

3


4

5

# MDP – Example



# MDP – Example




# MDP – Example


$S_t, A_t, R_t$  are random variables

Like dice rolls or coin flips, they can have different results for separate executions.

$$S_3 = s_{124}$$

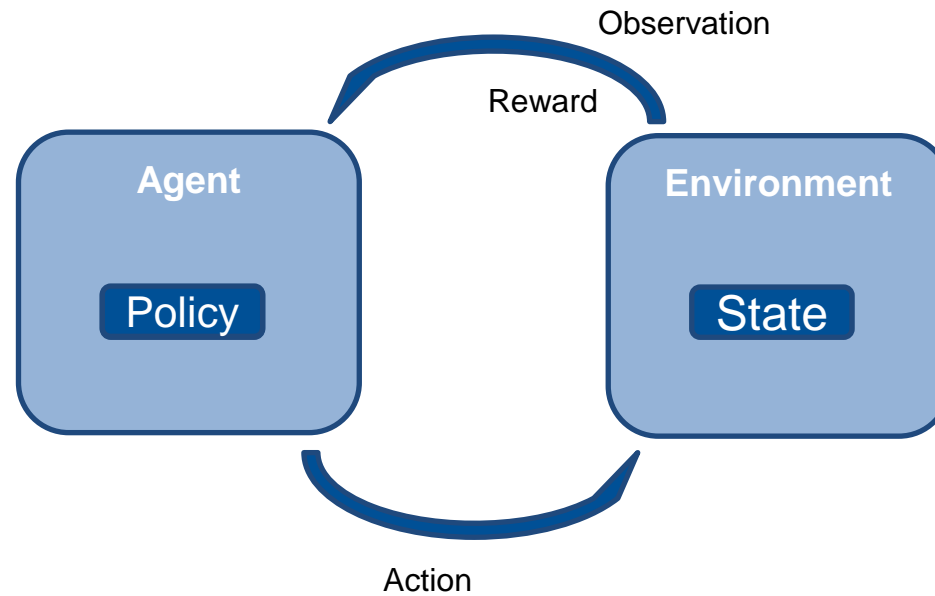


$$A_2 = a_2$$




$$\mathbb{P}(S_{t+1} = s_{124} | S_t = s_{75}, A_t = a_4) = ?$$

# MDP



# MDP – Definition



$\approx$

## Markov Decision Process

State Space:  $\mathcal{S} = \{s_1, s_2, \dots, s_N\}$

Action Space:  $\mathcal{A} = \{a_1, a_2, \dots, a_M\}$

Set of Rewards:  $\mathcal{R} \subset \mathbb{R}$

State-Transition

Probability Function:  $p: \mathcal{S} \times \mathcal{S} \times \mathcal{A} \rightarrow [0,1]$

Reward Function:  $R: \mathcal{R} \times \mathcal{S} \times \mathcal{A} \rightarrow [0,1]$

# MDP – Transition Function

State-Transition Probability Function:  $p: \mathcal{S} \times \mathcal{S} \times \mathcal{A} \rightarrow [0,1]$

$$p(s', s, a) := \mathbb{P}(S_{t+1} = s' | S_t = s, A_t = a)$$

- Defines how the environment behaves
- Function is usually not known
- Tic-tac-toe: contains rules of the games, as well as the opponents behavior

# MDP – Reward Function

Reward Function:  $R: \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow \mathbb{R}$

$$R(s, a, s')$$

- Defines what the environment rewards
- Function is usually designed by a human and therefore known
- Tic-tac-toe: contains rules of the games, as well as the opponents behavior

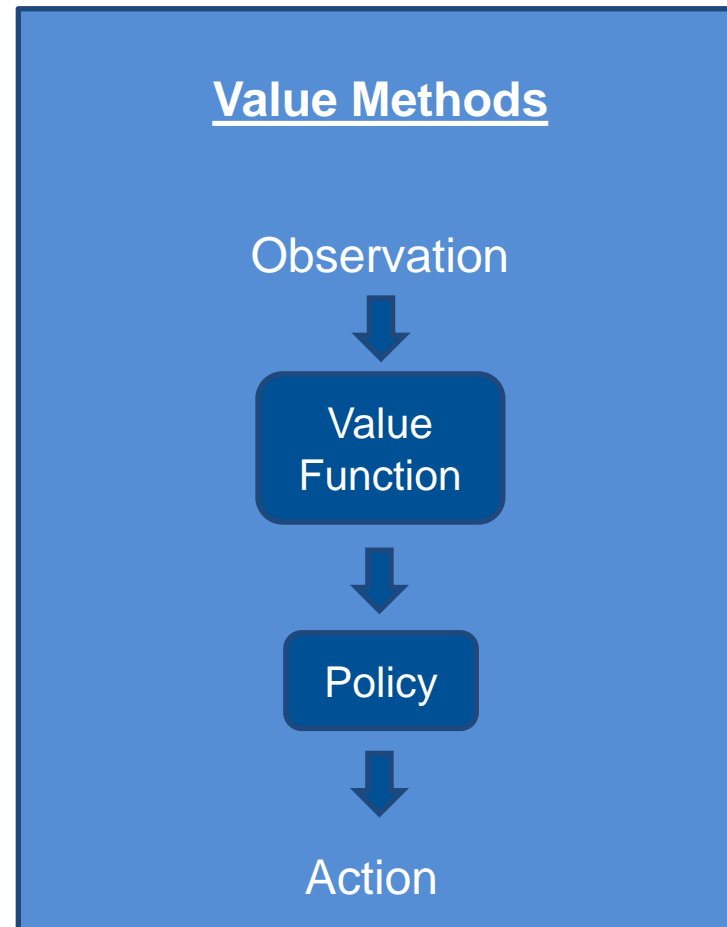


# MDP – Markov Property

The probability of each possible value of  $S_t$  and  $R_t$  depends **only** on the **immediately preceding** state and action and **not on earlier** states and actions.

„The future only depends on the present and not on the past“

$$\mathbb{P}(S_{t+1} = s' | S_t = s, A_t = a) = \\ \mathbb{P}(S_{t+1} = s' | S_t = s, A_t = a, \dots, S_0 = s'', A_0 = a'', )$$



# VF – Policy

Policy:  $\pi: \mathcal{S} \times \mathcal{A} \rightarrow [0,1]$

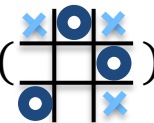
- defines how the agent behaves
- Rows are probability functions over the action space
- Example: random policy

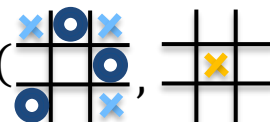
$\pi$	$a_1$	$a_2$	$a_3$	$a_4$
$s_1$	0.11	0.8	0.06	0.03
$s_2$	0.01	0.1	0.7	0.19
$s_3$	1.0	0.0	0.0	0.0
$s_4$	0.65	0.32	0.01	0.02
$s_5$	0.25	0.25	0.25	0.25
$s_6$	0.4	0.1	0.2	0.4

# VF – Value Functions

Value Functions:  $V: \mathcal{S} \rightarrow \mathbb{R}$   
 $Q: \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$


- indicate how „good“ a state or an action given a state is

$$V(\text{state}) = 0.64$$



$$Q(\text{state}, \text{action}) = 1.0$$


	$V$	$Q$	$a_1$	$a_2$	$a_3$	$a_4$
$s_1$	2.3	$s_1$	-0.1	4.3	0.3	1.8
$s_2$	-6.8	$s_2$	7.5	3	8.2	-5.7
$s_3$	10.0	$s_3$	0.0	10.0	0.0	0.0
$s_4$	-5.4	$s_4$	6.5	0.32	0.01	-0.2
$s_5$	4.1	$s_5$	3.0	0.21	-7.2	0.25
$s_6$	-0.4	$s_6$	4.5	0.1	-2.0	0.4


# Value Function




# Value Function



# Value Function




# Value Function






# Value Function




# Value Function

State:



Greedy Actions:



# VF – Motivation

You are:




Your opponent is:



It is your turn!

# VF – Motivation



# VF – Motivation


You are:



Your opponent is:




It is your turn!




indicates the next  
move

# VF – Motivation



# Which Value Function to choose?



Naive approach:

5 outcomes

3 positive

$$\Rightarrow V\left(\begin{array}{|c|c|c|} \hline \text{X} & \text{O} & \text{X} \\ \hline \text{O} & & \text{O} \\ \hline \text{O} & & \text{X} \\ \hline \end{array}\right) = \frac{3}{5} = 0.6$$


What if intermediate reward is given?

$\Rightarrow$  Need to add all reward until the end

$$G_t := \sum_{n=t+1}^T \gamma^n R_n$$

is also a random variable





Naive approach:

5 outcomes

3 positive

$$\Rightarrow V\left(\begin{array}{|c|c|c|} \hline \text{X} & \text{O} & \text{X} \\ \hline \text{O} & \text{X} & \text{O} \\ \hline \text{O} & \text{X} & \text{X} \\ \hline \end{array}\right) = \frac{3}{5} = 0.6$$

Implies equal probability for all outcomes.


What if agent only chooses actions from the second column if possible?

$$\Rightarrow V\left(\begin{array}{|c|c|c|} \hline \text{X} & \text{O} & \text{X} \\ \hline \text{O} & \text{X} & \text{O} \\ \hline \text{O} & \text{X} & \text{X} \\ \hline \end{array}\right) = \frac{2}{3} = 0.667$$

Value function depends on current policy (which we know)

$$\Rightarrow V_{\pi}(s)$$

# VF



Value function also depends on the environment (which we don't know)

=> Sampling

Sampling targets:

$$V_{\pi}(s) \approx \mathbb{E}_{\pi}[G_t | S_t = s]$$

$$Q_{\pi}(s, a) \approx \mathbb{E}_{\pi}[G_t | S_t = s, A_t = a]$$

# VF – Monte Carlo Method

1) Initialize with random policy and 0 for value function  $Q(s, a)$

$Q$	$a_1$	$a_2$	$a_3$	$a_4$
$s_1$	0.0	0.0	0.0	0.0
$s_2$	0.0	0.0	0.0	0.0
$s_3$	0.0	0.0	0.0	0.0
$s_4$	0.0	0.0	0.0	0.0
$s_5$	0.0	0.0	0.0	0.0
$s_6$	0.0	0.0	0.0	0.0


$\pi$	$a_1$	$a_2$	$a_3$	$a_4$
$s_1$	0.25	0.25	0.25	0.25
$s_2$	0.25	0.25	0.25	0.25
$s_3$	0.25	0.25	0.25	0.25
$s_4$	0.25	0.25	0.25	0.25
$s_5$	0.25	0.25	0.25	0.25
$s_6$	0.25	0.25	0.25	0.25

2) Play episode with current policy  $\pi$



# VF – Monte Carlo Method

2) Play episode with current policy  $\pi$



3) Iterate episode from back to front, calculate the return  $G_t$  and update the value function

4) Make policy  $\epsilon$ -greedy with regard to value function

5) continue with step 2)